

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

Workshop 1: Data Analysis with Pandas

```
In [2]: data = pd.read_csv('adult.csv')
```

```
In [3]: data.head()
```

Out[3]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class-label
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50K
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50K
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50K

Question no 1

```
In [4]: data.head(2)
```

Out[4]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class-label
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50K
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50K

This command shows the first two rows of the data set

```
In [5]: data.head(10)
```

Out[5]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50k
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50k
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50k
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50k
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50k
5	37	Private	284582	Masters	14	Married-civ-spouse	Exec-managerial	Wife	White	Female	0	0	40	United-States	<=50k
6	49	Private	160187	9th	5	Married-spouse-absent	Other-service	Not-in-family	Black	Female	0	0	16	Jamaica	<=50k
7	52	Self-emp-not-inc	209642	HS-grad	9	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	45	United-States	>50k
8	31	Private	45781	Masters	14	Never-married	Prof-specialty	Not-in-family	White	Female	14084	0	50	United-States	>50k
9	42	Private	159449	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	5178	0	40	United-States	>50k

this shows the first 10 rows of the data set

In [6]: data.tail(2)

Out[6]:

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	0	20	United-States	<=50k
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	0	40	United-States	>50k

This commands shows the last two rows of the data set

In [7]: data.shape

Out[7]: (32561, 15)

The data frame shape property tells you the dimensionalty of the data set in the form of number of rows and columns.

This data has 32561 rows and 15 columns.

Unique Data Set

In [39]: data = data.sample(n=30000, random_state = 70)

In [40]: data.shape

Out[40]: (30000, 14)

In [41]: data.describe()

Out[41]:

	age	education-num	capital-gain	capital-loss	hours-per-week
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	38.589133	10.080400	1086.949933	87.256033	40.43530
std	13.635182	2.572396	7459.713916	403.036258	12.34685
min	17.000000	1.000000	0.000000	0.000000	1.00000
25%	28.000000	9.000000	0.000000	0.000000	40.00000
50%	37.000000	10.000000	0.000000	0.000000	40.00000
75%	48.000000	12.000000	0.000000	0.000000	45.00000
max	90.000000	16.000000	99999.000000	4356.000000	99.00000

This command gives the description of the data. It shows the mean, standard deviation, count, minimum value maximum value and percentiles.

In [11]: data['education-num'].value_counts()

Out[11]:

9	9676
10	6714
13	4955
14	1578
11	1267
7	1081
12	984
6	861
4	598
15	524
5	483
8	393
16	383
3	304
2	153
1	46

Name: education-num, dtype: int64

In [12]: data = data.drop(['fnlwgt'], axis=1)

This command drops the column 'fnlwgt' and I have used axis=1 to drop the first row.

In [13]: data.shape

Out[13]: (30000, 14)

As I have dropped one coulumn 'fnlwgt' from the data so now the data has 14 columns.

In [14]: data.describe(include='all')

Out[14]:

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week
count	30000.000000	30000	30000	30000.000000	30000	30000	30000	30000	30000	30000.000000	30000.000000	30000.000000
unique	NaN	9	16	NaN	7	15	6	5	2	NaN	NaN	NaN
top	NaN	Private	HS-grad	NaN	Married-civ-spouse	Prof-specialty	Husband	White	Male	NaN	NaN	NaN
freq	NaN	20927	9676	NaN	13791	3818	12166	25615	20049	NaN	NaN	NaN
mean	38.589133	NaN	NaN	10.080400	NaN	NaN	NaN	NaN	NaN	1086.949933	87.256033	40.43530
std	13.635182	NaN	NaN	2.572396	NaN	NaN	NaN	NaN	NaN	7459.713916	403.036258	12.34685
min	17.000000	NaN	NaN	1.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000	1.00000
25%	28.000000	NaN	NaN	9.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000	40.00000
50%	37.000000	NaN	NaN	10.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000	40.00000
75%	48.000000	NaN	NaN	12.000000	NaN	NaN	NaN	NaN	NaN	0.000000	0.000000	45.00000
max	90.000000	NaN	NaN	16.000000	NaN	NaN	NaN	NaN	NaN	99999.000000	4356.000000	99.00000

This commands shows the description of all the variables we have in data set.

In [15]: data['education'].nunique()

Out[15]: 16


This commands tell about the unique values we have in education column.

```
In [16]: data['age'].value_counts()
```

```
Out[16]: 31    824
35    821
36    812
23    807
28    804
...
83      5
85      3
88      3
86      1
87      1
Name: age, Length: 73, dtype: int64
```


```
In [17]: data.boxplot(column='age')
```

```
Out[17]: <AxesSubplot:>
```




```
In [18]: data['age'].hist(bins=100)
```

```
Out[18]: <AxesSubplot:>
```



```
In [19]: data.age.hist(bins=100)
```

```
Out[19]: <AxesSubplot:>
```



```
In [20]: data['sex'].value_counts()
```

```
Out[20]: Male      20049
Female    9951
Name: sex, dtype: int64
```

The data is collected from 20049 males and 9951 females.

```
In [21]: data.columns
```

```
Out[21]: Index(['age', 'workclass', 'education', 'education-num', 'marital-status',
        'occupation', 'relationship', 'race', 'sex', 'capital-gain',
        'capital-loss', 'hours-per-week', 'native-country', 'class-label'],
        dtype='object')
```

```
In [22]: data['workclass'].value_counts()
```

```
Out[22]: Private      20927
Self-emp-not-inc    2336
Local-gov          1925
?                  1694
State-gov          1191
Self-emp-inc       1020
Federal-gov        888
Without-pay        13
Never-worked        6
Name: workclass, dtype: int64
```

Question no 2

```
In [23]: data['sex'].value_counts()
```

```
Out[23]: Male      20049
Female    9951
Name: sex, dtype: int64
```

Applying groupby functions in order to summarise the data.

Groupby functions are usually used with aggregate functions, which are useful to summarise the dataset and make observations. Some common functions are SUM, MEAN, MAX, MIN and COUNT. Using groupby, we can answer questions such as:

Question: What is the average age of each gender in the given population?

```
In [24]: data['age'].groupby([data['sex']]).mean()
```

```
Out[24]: sex
Female    36.851271
Male      39.451693
Name: age, dtype: float64
```

This shows that average age of female is 36 and average age of male is 39 in the adult data.

Question. What is the average age of male and female across different education categories?

```
In [25]: data['age'].groupby([data['sex'], data['education']]).mean()
```

```
Out[25]: sex      education
Female  10th      35.319703
        11th      30.348148
        12th      30.150376
        1st-4th   49.976190
        5th-6th   45.285714
        7th-8th   50.165563
        9th       41.789855
        Assoc-acdm 36.413265
        Assoc-voc  37.823276
        Bachelors  35.619906
        Doctorate  45.120482
        HS-grad    38.593172
        Masters    42.932515
        Preschool  42.266667
        Prof-school 40.716049
        Some-college 33.788454
Male    10th      38.094595
        11th      33.331361
        12th      32.826923
        1st-4th   45.684685
        5th-6th   41.656388
        7th-8th   48.255034
        9th       40.492754
        Assoc-acdm 38.064189
        Assoc-voc  39.022416
        Bachelors  40.395604
        Doctorate  48.160000
        HS-grad    39.178997
        Masters    44.525253
        Preschool  42.322581
        Prof-school 45.584650
        Some-college 37.012582
Name: age, dtype: float64
```

In the above code, we group by 'sex' and 'education' and computed mean 'age' in the given population.

NOTE: groupby can be applied to multiple columns.

NOTE: groupby can be applied on numeric attributes only.

Question no 3

What is the average contribution to capital-gain of each sex and occupation category?

```
In [28]: #Answer
data['capital-gain'].groupby([data['sex'],data['occupation']]).mean()
```

```
Out[28]: sex      occupation      capital-gain
Female   ?              351.420716
         Adm-clerical    508.543497
         Craft-repair    807.793269
         Exec-managerial 1022.757263
         Farming-fishing 1293.019231
         Handlers-cleaners 151.421769
         Machine-op-inspct 149.511583
         Other-service    160.582691
         Priv-house-serv  302.651163
         Prof-specialty   1304.731568
         Protective-serv  1734.301370
         Sales            281.543199
         Tech-support     658.773292
         Transport-moving 455.589744
Male     ?              877.041394
         Adm-clerical    480.800352
         Armed-Forces     0.000000
         Craft-repair    659.414846
         Exec-managerial 2778.056962
         Farming-fishing 504.397390
         Handlers-cleaners 286.047748
         Machine-op-inspct 397.674191
         Other-service    253.938672
         Priv-house-serv  74.250000
         Prof-specialty   3485.083850
         Protective-serv  606.676864
         Sales            1951.053906
         Tech-support     724.552876
         Transport-moving 494.525706
Name: capital-gain, dtype: float64
```

In the above code, we group by 'sex' and 'occupation' and computed mean 'capital-gain' in the given population

Question no 4

Identify the average capital-gain by males and females accross different marital-status.

```
In [30]: #Answer
data['capital-gain'].groupby([data['sex'],data['marital-status']]).mean()
```

```
Out[30]: sex      marital-status      capital-gain
Female   Divorced          454.577590
         Married-AF-spouse  204.076923
         Married-civ-spouse 1615.607662
         Married-spouse-absent 373.540404
         Never-married      335.807964
         Separated          366.775891
         Widowed            493.536137
Male     Divorced          1157.684535
         Married-AF-spouse  810.888889
         Married-civ-spouse 1791.060031
         Married-spouse-absent 1037.455026
         Never-married      434.198822
         Separated          872.103825
         Widowed            925.869281
Name: capital-gain, dtype: float64
```

In the above code, we group by 'sex' and 'marital-status' and computed mean 'capital-gain' in the given population

Question. What is the maximum age accross differnt races?

Let's first see what are the different races and then apply groupby.

```
In [32]: data['race'].value_counts()
```

```
Out[32]: White          25615
Black          2890
Asian-Pac-Islander  961
Amer-Indian-Eskimo  281
Other          253
Name: race, dtype: int64
```

```
In [33]: data['age'].groupby([data['race']]).max()
```

```
Out[33]: race
Amer-Indian-Eskimo    82
Asian-Pac-Islander    90
Black                 90
Other                 77
White                 90
Name: age, dtype: int64
```

It reflects that maximum adult of age 82 is amer-indian-eskimo Maximum age of Asian-Pac-islander in the data is 90 Maximum age of Black person in the data is 90 Maximum age of White person in the data is 90

Question no 5

Are minimum and maximum age by sex same?

Minimum age by sex

```
In [35]: data['age'].groupby([data['sex']]).min()
```

```
Out[35]: sex
Female    17
Male      17
Name: age, dtype: int64
```

```
In [36]: data['age'].groupby([data['sex']]).max()
```

```
Out[36]: sex
Female    90
Male      90
Name: age, dtype: int64
```

Yes, the minimum and maximum age by sex is same

Data Visualisation

Matplotlib is python library for visualising data in the form of graphs such as histograms, scatter, box plot, line plots, heat plots, etc.

```
In [37]: import matplotlib.pyplot as plt
%matplotlib inline
```


```
In [42]: data.describe()
```

```
Out[42]:
```

	age	education-num	capital-gain	capital-loss	hours-per-week
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	38.589133	10.080400	1086.949933	87.256033	40.43530
std	13.635182	2.572396	7459.713916	403.036258	12.34685
min	17.000000	1.000000	0.000000	0.000000	1.00000
25%	28.000000	9.000000	0.000000	0.000000	40.00000
50%	37.000000	10.000000	0.000000	0.000000	40.00000
75%	48.000000	12.000000	0.000000	0.000000	45.00000
max	90.000000	16.000000	99999.000000	4356.000000	99.00000

```
In [43]: data['age'].hist(bins=50)
```

```
Out[43]: <AxesSubplot:>
```




Histograms are used to represent the distribution of dataset. The bars of the histograms are known as bins or "buckets" - the range of

histograms is used to represent the distribution of dataset. The bars of the histograms are known as bins or bucket – the range of values. Bins are of same width. Width of the bins can be calculated as (max value of data – min value of data) / total number of bins. The bins are usually specified as continuous, non-overlapping intervals of a variable.

In the above figure, histogram with bins = 50 is used to show number of people belongs to different age-groups. Here, x-axis represents 'age' and y-axis represents the 'count'. **Try-it-yourself:** change bins = 100 and run the cell to observe the difference for your own understanding.

```
In [46]: data['age'].hist(bins=100)
```


```
Out[46]: <AxesSubplot:>
```



if we increase the bin size, the grouping in histogram is more clearly visible

```
In [48]: data.boxplot(column='age')
```


```
Out[48]: <AxesSubplot:>
```



In the above figure, boxplot is used to find the average number of people belongs to which age-range group. The mean is around 38 age. and there are outliers after 78 age. the minimum age we can see from box plot is 17 and maximum age is 78. After 78 age there are outliers.


```
In [49]: data['capital-gain'].hist(bins=100)
```

```
Out[49]: <AxesSubplot:>
```




```
In [50]: data.boxplot(column='capital-gain')
```

```
Out[50]: <AxesSubplot:>
```

```
In [51]: data.boxplot(column='age', by = 'education', grid=False, rot = 45, fontsize = 10)
```

```
Out[51]: <AxesSubplot:title={'center':'age'}, xlabel='education'>
```

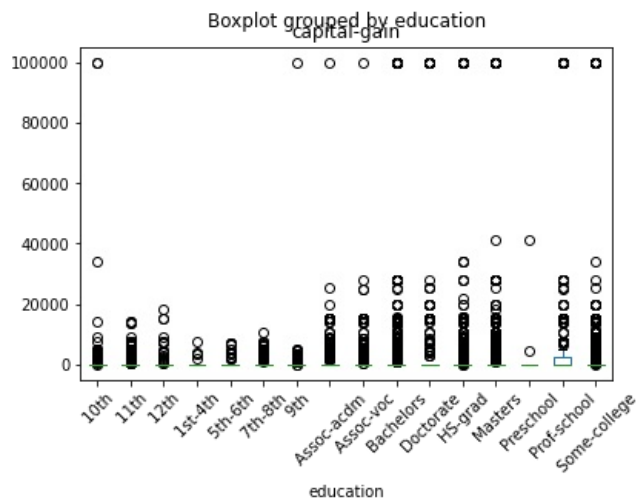


```
In [52]: data['education'].value_counts()
```

```
Out[52]: HS-grad      9676
Some-college  6714
Bachelors    4955
Masters      1578
Assoc-voc    1267
11th         1081
Assoc-acdm   984
10th         861
7th-8th      598
Prof-school  524
9th          483
12th         393
Doctorate    383
5th-6th      304
1st-4th      153
Preschool    46
Name: education, dtype: int64
```

```
In [53]: data.boxplot(column='capital-gain', by = 'education', grid=False, rot = 45, fontsize = 10)
```

```
Out[53]: <AxesSubplot:title={'center':'capital-gain'}, xlabel='education'>
```



After performing some basic data analysis, let's look at data pre-processing to improve the quality of the dataset.

Data pre-processing is an important step in the process. Raw data can be unstructured and full of noise. Aim of this phase is to clean the raw data, reduce noise and to prepare the dataset that can be accepted by the algorithm as an input. Remember garbage in, garbage out!

```
In [54]: data['marital-status'].value_counts()
```

```
Out[54]: Married-civ-spouse      13791
Never-married                  9827
Divorced                       4104
Separated                      955
Widowed                        914
Married-spouse-absent          387
Married-AF-spouse               22
Name: marital-status, dtype: int64
```

Checking NULL values in the dataset

```
In [55]: data.apply(lambda x: sum(x.isnull()), axis = 0)
```

```
Out[55]: age                0
workclass                 0
education                 0
education-num             0
marital-status            0
occupation                0
relationship              0
race                     0
sex                      0
capital-gain              0
capital-loss              0
hours-per-week            0
native-country            0
class-label              0
dtype: int64
```

As the missing values in this data is already replaces by ?.

Data Transformation

Label encoding:

Some attributes are categorical, therefore (statistical) analysis on those variables is not possible. We need to convert all categorical variables (string labels) into numeric by encoding the categories. Package 'sklearn' provides 'LabelEncoder' library for encoding labels

between 0 to n-1 discrete values/labels, where n is the number of values/labels. E.g.: Male -> 0 Female -> 1

```
In [56]: from sklearn.preprocessing import LabelEncoder
```

```
In [57]: data.head()
```

```
Out[57]:
```

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class-label
31113	28	Private	Assoc-voc	11	Married-civ-spouse	Prof-specialty	Wife	White	Female	0	0	5	United-States	<=50K
12788	24	State-gov	Doctorate	16	Never-married	Prof-specialty	Not-in-family	White	Female	0	0	99	England	<=50K
27524	38	Private	HS-grad	9	Separated	Sales	Not-in-family	White	Male	0	0	60	United-States	<=50K
30497	39	Self-emp-not-inc	10th	6	Married-spouse-absent	Other-service	Not-in-family	White	Female	0	1721	15	United-States	<=50K
9118	23	Private	10th	6	Never-married	Handlers-cleaners	Other-relative	Other	Male	0	0	40	United-States	<=50K

```
In [58]: data.dtypes
```

```
Out[58]: age                int64
workclass                object
education                object
education-num            int64
marital-status           object
occupation               object
relationship             object
race                    object
sex                     object
capital-gain             int64
capital-loss             int64
hours-per-week           int64
native-country           object
class-label              object
dtype: object
```

```
In [59]: columns = list(data.select_dtypes(exclude=['int64']))
```

As we do not need to convert the integers they are already in numbers. So, I drop all the integer columns in the data.

```
In [60]: columns
```

```
Out[60]: ['workclass',
'education',
'marital-status',
'occupation',
'relationship',
'race',
'sex',
'native-country',
'class-label']
```

```
In [61]: data['class-label'].value_counts()
```

```
Out[61]: <=50K    22768
>50K        7232
Name: class-label, dtype: int64
```

```
In [108]: le = LabelEncoder()
for i in columns:
    #print(i)
    data[i] = le.fit_transform(data[i])
data.dtypes
```

```
Out[108]: age                int64
workclass                int64
education                int64
education-num            int64
marital-status           int64
occupation               int64
relationship             int64
race                    int64
sex                     int64
capital-gain             int64
capital-loss             int64
hours-per-week           int64
native-country           int64
class-label              int64
dtype: object
```

```
In [109]: data.head()
```

```
Out[109]:
```

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class-label
31113	28	4	14	11	2	2	5	4	0	0	0	5	33	0
12788	24	7	2	16	4	2	1	4	0	0	0	99	41	0
27524	38	4	3	9	5	4	1	4	1	0	0	60	33	0
30497	39	6	0	6	3	13	1	4	0	0	1721	15	33	0
9118	23	4	0	6	4	11	2	3	1	0	0	40	33	0

```
In [110]: data['workclass'].value_counts()
```


```
Out[110]:
```

4	20927
6	2336
2	1925
0	1694
7	1191
5	1020
1	888
8	13
3	6

Name: workclass, dtype: int64

You will notice that all the values are now numeric. Now, more computations and analysis can be performed on the dataset.

```
In [111]: data['education'].hist(bins=100)
```



```
In [112]: data.describe(include='all')
```

```
Out[112]:
```

	age	workclass	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	class-label
count	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000	30000.000000
mean	38.589133	3.867367	7.078300	10.080400	2.611067	6.341533	1.445167	3.665333	0.668300	10.000000	0.000000	40.000000	33.000000	0.000000
std	13.635182	1.455648	4.831344	2.572396	1.507229	4.263875	1.605411	0.848502	0.470832	74.000000	0.000000	11.000000	33.000000	0.000000
min	17.000000	0.000000	0.000000	1.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000
25%	28.000000	4.000000	3.000000	9.000000	2.000000	2.000000	0.000000	4.000000	0.000000	0.000000	0.000000	40.000000	33.000000	0.000000
50%	37.000000	4.000000	7.000000	10.000000	2.000000	6.000000	1.000000	4.000000	1.000000	0.000000	0.000000	40.000000	33.000000	0.000000
75%	48.000000	4.000000	12.000000	12.000000	4.000000	9.000000	3.000000	4.000000	1.000000	0.000000	0.000000	40.000000	33.000000	0.000000
max	90.000000	8.000000	15.000000	16.000000	6.000000	14.000000	5.000000	4.000000	1.000000	999.000000	0.000000	99.000000	41.000000	0.000000

Report

Question no 6

Write a summary of the outcome of `data.describe()`

Answer When we describe the data, it will give the total number of values in each column. It also give the mean of each variable of the data. Like in **age** the mean is 38. It represents that avarage age is 38 in the data. The describe function also gives the maximum and minimum value and standard deviation of each variable. The minimum age is 17 and maximum is 90. it also gives 25% percentile, 50 percentile and 75 percentile of the data.

Question no 7

What are the different data types (or attribut types) in data mining? Explain with the help of the examples from Adult dataset. HINT: Don't get confused with data types in Python or Pandas.

Answer There are mainly two attributes in the data mining

- Quantitative attribute such as discrete and conitnuous attribute
- Qualitative attribute such as oridnal, nominal and binary attributes

Question no 8

```
In [96]: data1 = pd.read_csv('adult.csv')
```

Highest migrants belongs to which country?

```
In [97]: data1['native-country'].value_counts()
```

```
Out[97]: United-States      29170
Mexico          643
?               583
Philippines     198
Germany         137
Canada          121
Puerto-Rico    114
El-Salvador     106
India           100
Cuba            95
England         90
Jamaica         81
South           80
China           75
Italy           73
Dominican-Republic 70
Vietnam         67
Guatemala       64
Japan           62
Poland          60
Columbia        59
Taiwan          51
Haiti           44
Iran            43
Portugal        37
Nicaragua       34
Peru            31
France          29
Greece          29
Ecuador         28
Ireland         24
Hong            20
Cambodia        19
Trinidad&Tobago 19
Laos            18
Thailand        18
Yugoslavia      16
Outlying-US(Guam-USVI-etc) 14
Honduras        13
Hungary         13
Scotland        12
Holand-Netherlands 1
Name: native-country, dtype: int64
```

39 is assigned to United States. Most adults are from United States in the data.

Question no 9

Which occupation represents more males than females?

```
In [103]: data1['sex'].groupby(data1['occupation']).value_counts()
```

```
Out[103]: occupation      sex      count
?                Male      1002
                Female     841
Adm-clerical      Female     2537
                Male      1233
Armed-Forces      Male        9
Craft-repair      Male     3877
                Female     222
Exec-managerial   Male     2907
                Female     1159
Farming-fishing   Male      929
                Female      65
Handlers-cleaners Male     1206
                Female     164
Machine-op-inspct Male     1452
                Female     550
Other-service     Female     1800
                Male     1495
Priv-house-serv   Female     141
                Male        8
Prof-specialty    Male     2625
                Female     1515
Protective-serv   Male      573
                Female      76
Sales             Male     2387
                Female     1263
Tech-support      Male      580
                Female     348
Transport-moving  Male     1507
                Female      90
Name: sex, dtype: int64
```

Almost all the occupation has more males than females, except adm-clerical, other service and pric-house-serv.

Question no 10

What is the difference between data.head() and data.tail()?

Answer Data.head shows the first 5 rows of the dataframe However, Data,tail() shows that last 5 rows of the data set

```
In [105]: data1.head()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	clas
0	39	State-gov	77516	Bachelors	13	Never-married	Adm-clerical	Not-in-family	White	Male	2174	0	40	United-States	<=50
1	50	Self-emp-not-inc	83311	Bachelors	13	Married-civ-spouse	Exec-managerial	Husband	White	Male	0	0	13	United-States	<=50
2	38	Private	215646	HS-grad	9	Divorced	Handlers-cleaners	Not-in-family	White	Male	0	0	40	United-States	<=50
3	53	Private	234721	11th	7	Married-civ-spouse	Handlers-cleaners	Husband	Black	Male	0	0	40	United-States	<=50
4	28	Private	338409	Bachelors	13	Married-civ-spouse	Prof-specialty	Wife	Black	Female	0	0	40	Cuba	<=50

```
In [106]: data1.tail()
```

	age	workclass	fnlwgt	education	education-num	marital-status	occupation	relationship	race	sex	capital-gain	capital-loss	hours-per-week	native-country	
32556	27	Private	257302	Assoc-acdm	12	Married-civ-spouse	Tech-support	Wife	White	Female	0	0	38	United-States	
32557	40	Private	154374	HS-grad	9	Married-civ-spouse	Machine-op-inspct	Husband	White	Male	0	0	40	United-States	
32558	58	Private	151910	HS-grad	9	Widowed	Adm-clerical	Unmarried	White	Female	0	0	40	United-States	
32559	22	Private	201490	HS-grad	9	Never-married	Adm-clerical	Own-child	White	Male	0	0	20	United-States	
32560	52	Self-emp-inc	287927	HS-grad	9	Married-civ-spouse	Exec-managerial	Wife	White	Female	15024	0	40	United-States	

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js